

Parallax Bumpmapping

Whitepaper

Copyright © Imagination Technologies Ltd. All Rights Reserved.

This publication contains proprietary information which is subject to change without notice and is supplied 'as is' without warranty of any kind. Imagination Technologies and the Imagination Technologies logo are trademarks or registered trademarks of Imagination Technologies Limited. All other logos, products, trademarks and registered trademarks are the property of their respective owners.

Filename : Parallax Bumpmapping.Whitepaper.1.0.5.External.doc
Version : 1.0.5 External Issue (Package: PowerVR SDK REL_3.0@2149525)
Issue Date : 30 Aug 2012
Author : Imagination Technologies Ltd

Contents

1. Introduction	3
2. Description	4
2.1. Overview	4
2.2. Required Information	4
2.3. How It Works	5
3. Optimizations	6
3.1. Initial Prototyping	6
3.2. On Platform.....	6
4. Analysis	8
4.1. Advantages.....	8
4.2. Disadvantages	8
5. Conclusion	9
6. Bibliography	10
7. Contact Details.....	11

List of Figures

Figure 2-1 Observed Texel vs. Real Texel	4
Figure 2-2 New UV Coordinate	5
Figure 4-1 Texture Crawling (No Scaling).....	8
Figure 4-2 Texture Crawling (Scaled)	8

1. Introduction

The parallax effect is exhibited when areas of a surface appear to move relative to one another as the position of the viewer changes. This is visible on any non-flat surface, a stone wall, or a speaker stack for example. It is easy to simulate parallax within a scene using geometry as the mathematics used to display a three dimensional world in a computer, generate parallax as a by-product of the display process. This is, however, computationally expensive to add to every rough or lined surface due to the additional geometry required.

The method described here provides a means to approximate the parallax effect using height data and the view position to adjust the UV coordinates of an object; all without the addition of further polygons. In addition, optimizations will be discussed for a parallax bump mapping shader specifically optimized for PowerVR SGX platforms.

2. Description

2.1. Overview

On a real surface the variations in height produce a parallax effect; with non-parallax bump mapping there are no height variations on a polygon.




Figure 2-1 Observed Texel vs. Real Texel

As can be seen in Figure 2-1 the observed Texel T_b corresponds to point A when using a flat polygon. If the object were real, the eye would see Texel T_b as this corresponds with the point of intersection of the 'real surface' with the view vector. The goal of parallax bump mapping is to adjust the UV coordinates used at point A so they match those of point B; if this is done per pixel over an entire surface then a parallax effect should emerge.

2.2. Required Information

Reasonably advanced graphics systems often describe objects with a diffuse colour, a normal map, and occasionally a specular map. In addition to this information, the parallax effect requires the use of a height map. This height map stores information pertaining to the height of a surface at a given point on it. Finally, the view direction is required to ensure that the texture coordinate is moved in the correct direction.

As the calculations are being used to determine a modified UV coordinate they will be done in tangent space; as such there will be a requirement for a world to tangent space matrix. This matrix will be used to transform the view direction into tangent space on a per vertex basis (per pixel information is determined by interpolating this data across a polygon).

2.3. How It Works

In standard texture mapping each vertex has a UV within the texture: when each polygon is rendered this UV is interpolated across the surface to determine the segment of the texture that is to be sampled. In the case of parallax, instead of sampling the texture the coordinates are used to sample a height map. With this information, a ray is traced from point A until it intersects the view vector; the coordinate of this new point is then used to sample all of the remaining maps (texture, normal etc.).




Figure 2-2 New UV Coordinate

As can be seen in Figure 2-2, the result is only an approximation; techniques do exist for more accurate simulations, but in all cases they are far more computationally expensive. An example of one of these techniques would be Parallax Occlusion Bump Mapping with Reverse Height Tracing (Brawley & Tatarchuk, 2004).

The offset $\overline{AB_{new}}$ is calculated by first converting the view direction (V) into tangent space and normalizing it. The normalized view direction contains x and y components that lie on the plane of the surface and a z component which is perpendicular to the surface. Next the height value (in the range 0.0 to 1.0) is scaled and biased and then multiplied by the x and y components of the view direction; the resulting vector is finally added to the original UV coordinate to give the offset UV.

The original equation for determining the texture offset is:

$$T_{new} = T_a + \frac{V_{(x,y)} * h_{scaled}}{V_z}$$

Using this equation V_z begins to cause problems at grazing angles where it begins to become close to zero; the offset values become unusably high, breaking the parallax effect.

To solve this V_z is dropped from the equation entirely, limiting the offset to no more than the value of h_{scaled} . At grazing angles this improves the image quality with little difference to the overall quality in general. This gives a final equation of:

$$T_{new} = T_a + (V_{(x,y)} * h_{scaled})$$

3. Optimizations

Parallax bump mapping, while considerably less expensive computationally than other available options, is still quite expensive for mobile platforms; as such, much care was taken to optimize the effect.

3.1. Initial Prototyping

Initial development was done in PVRShaman with all precisions set to 'highp'. The counts from PVRShaman's editor window were used to suggest areas for improvement. The primary areas identified were matrix multiplication within the vertex shader, and normalisation within the fragment shader. As the first identified optimization, a faster means of creating an eye to tangent space conversion matrix was devised. This optimization used the inverse transpose of the world view matrix to create a tangent, Binormal and normal; these were then used to manually create the eye to tangent space matrix.

```
highp vec3 n = normalize(worldViewIT * vertNormal);  
highp vec3 t = normalize(worldViewIT * vertTangent);  
highp vec3 b = cross(n,t);  
  
// Create the matrix from the above  
mat3 mEyeToTangent = mat3( t.x, b.x, n.x,  
                           t.y, b.y, n.y,  
                           t.z, b.z, n.z);
```

This moved some of the grunt work to the CPU (specifically calculating worldViewIT which would be passed as a uniform). On investigation little more could be gained from optimizing the vertex shader without changing precisions.

The next identified optimization was in the various normalizes within the fragment shader. The fragment shader for parallax bump mapping usually begins by normalizing the light and view directions. It was possible to move the normalization of the view direction into the vertex shader with a very minor drop in visual quality. While this did increase the cycle count of the vertex shader it reduced the cycle count of the fragment shader by a commensurate amount; this proved to be a small gain as the fragment shader is run far more often than the vertex shader. This was attempted with the light direction, but the drop in visual quality was significant, and deemed unacceptable.

3.2. On Platform

Once the initial prototyping had been completed the effect was tested on a device using PVRTune to profile the applications performance. Initial results showed the effect was bandwidth limited. Parallax bump mapping requires at least three texture reads: a height map, a normal map and a texture map. Of these, only the height map is an independent read, the other two being dependent on the initial read of an offset calculation. As such, only the height map read could be prescheduled. To amortise this cost the textures were all compressed using PVRTC 2 bpp (bits per pixel); this provided a considerable gain in performance, approximately 20 fps (frames per second).

With the textures now heavily compressed PVRTune was once again used to profile the application. This second run identified that the effect had become fragment shader limited so this became the area of focus for the next pass of optimizations. The first fragment shader optimization was to move the multiplication by fParallaxScale into the vertex shader. As with the normalization of the view direction this provides a gain as there are more fragments being processed than vertices. This saved a single cycle per fragment. While this might seem a small gain as compared to the texture compression, these small gains can add up over many profiling and optimisation passes to a significant improvement. It is important to note that each stage further profiling was performed, without repeated profiling time would have been wasted optimising areas of the pipeline that were not performance limiting, or on fixing bottlenecks that had moved.

Attention was then returned to the dependent texture reads. When a dependent texture read occurs the current thread running on the USSE within the SGX is suspended until the comparatively slow read is completed, and another thread is swapped in. By re-arranging the code within the shader so that the two dependent texture reads were back-to-back, another performance improvement was gained. This works because the compiler can now batch these dependent texture reads together, whereas previously the shader looked as follows:

- Mathematics
- Texture Read
- Wait
- Mathematics
- Texture Read
- Wait

It now looks as follows:

- Mathematics
- Texture Read
- Texture Read
- Wait

This effectively hides some of the wait time of the first texture read behind the execution time of the second texture read. Unfortunately a wait of some description is inevitable due to the limitations of dependent texture reads. Nonetheless, the strong parallelism of the PowerVR SGX hardware ensures that these waits can be used effectively by other threads (see [SGX Architecture Guide for Developers](#)).

The last and most difficult area to optimize was the precision settings. Initially, everything used high precision (highp), however there are additional costs to highp; where possible low precision (lowp) should be used. After some testing, it was determined that all calculations done within the fragment shader could be lowp with little decrease in visual appeal. Likewise, the view direction could be lowp as could the texture coordinate (in both cases, because all values were in the range of 0-1 only precision was lost, not actual integer information). Light direction was also moved to lowp; initially this damaged the overall effect, however, moving the lowp to highp conversion into the vertex shader provided a gain of 4 cycles in the fragment shader at a cost of 2 cycles in the vertex shader.

The key to producing gains from the use of lowp is to be aware of the cost of converting from one precision to another; if a calculation includes a highp value then there will be no performance gains from using a lowp value in the same calculation.

With the above performance enhancements completed, the shader ran at ~80fps on the non-vsync limited test device.

4. Analysis

4.1. Advantages

While this parallax bump mapping effect is still computationally expensive on mobile devices the alternatives, greater model complexity, displacement mapping etc. are far more expensive. Parallax bump mapping provides an adequate solution that is artistically pleasing using simple geometry at a reduced cost to all the alternatives.

4.2. Disadvantages

There are several disadvantages to parallax bump mapping; the first is the lack of a complex silhouette. While parallax bump mapping allows for the simulation of false geometry within a polygon, the silhouette of that polygon does not change. This can be a problem on surfaces with steep complex bump maps when viewed at extreme angles. Effectively, the lack of a complex silhouette breaks the visual cues that allow the parallax effect to function.

Another disadvantage is a lack of self-shadowing; areas of the height map that are higher than other others do not cast shadows over the rest of the height map. Techniques exist that can simulate this via reverse height tracing; however, at the time of writing it was decided that the computational cost of these techniques was too expensive for widespread use in the current generation of mobile graphics hardware.

Finally, the parallax technique makes an assumption that, as can be seen in Figure 2-2 does not always hold true, that point **A** and point **B_{new}** both lie on the surface of the 'real surface'. The effect of this assumption is that artifacting occurs at points of steep height change. This becomes even more pronounced at points with vertical or near vertical height changes where the texture can be seen to crawl up the sides of the raised section. In the case of this demo the effect has been minimized through the use of a scaling vector that restricts how much the UV can be perturbed.




Figure 4-1 Texture Crawling (No Scaling)



Figure 4-2 Texture Crawling (Scaled)

5. Conclusion

Relative to other techniques that perform similar functions the parallax bump mapping technique is an inexpensive means to better simulate complex surfaces than bump mapping alone. The use of bump mapping helps to offset some of the issues with parallax mapping by itself, specifically the lack of lighting and shadowing, though not self-shadowing. Parallax mapping provides additional visual cues that bump mapping alone does not, that help improve the overall effect.

With correct optimizations the effect can be used widely on PowerVR SGX platforms reaching acceptable performance levels even when running full screen on lower end devices. Future PowerVR platforms should be more than capable of running full scene parallax bump mapping, possibly including some more advanced techniques for complex silhouettes and self-shadowing that are impractical presently.

6. Bibliography

Brawley, Z., & Tatarchuk, N. (2004). "Parallax Occlusion Mapping: Self-Shadowing, Perspective-Correct Bump Mapping Using Reverse Height Map Tracing. In *Shader X3*.

7. Contact Details

For further support contact:

devtech@imgtec.com

Developer Technology
Imagination Technologies Ltd.
Home Park Estate
Kings Langley
Herts, WD4 8LZ
United Kingdom

Tel: +44 (0) 1923 260511

Fax: +44 (0) 1923 277463

Alternatively, you can use the PowerVR Insider forums:

www.imgtec.com/forum

For more information about PowerVR or Imagination Technologies Ltd. visit our web pages at:

www.imgtec.com

Imagination Technologies, the Imagination Technologies logo, AMA, Codescape, Enigma, IMGworks, I2P, PowerVR, PURE, PURE Digital, MeOS, Meta, MBX, MTX, PDP, SGX, UCC, USSE, VXD and VXE are trademarks or registered trademarks of Imagination Technologies Limited. All other logos, products, trademarks and registered trademarks are the property of their respective owners.